

Reproducing Grounding DINO: Open-Vocabulary Object Detection and Visual Grounding on COCO 2017

Su Siqi (12310645), Yang Guanyuhan (12313614), Mo Fengyuan (12311805)
Southern University of Science and Technology (SUSTech), Shenzhen, China

https://github.com/YangGuanyuhan/CV_Project_2026_S

Abstract—Open-vocabulary object detection extends traditional detection to arbitrary text-described categories, enabling models to detect objects beyond a fixed training vocabulary. This report presents the reproduction and fine-tuning of Grounding DINO, a state-of-the-art open-vocabulary detector, on the COCO 2017 dataset. Using the official Swin-Tiny backbone checkpoint, we first evaluate the zero-shot baseline, achieving an mAP of 0.485 on the COCO val2017 set. We then fine-tune the model for 6 epochs on the COCO train2017 split, saving a checkpoint after each epoch. The best performing checkpoint (epoch 5) achieves an mAP of 0.559, an AP50 of 0.735, and an AP75 of 0.615, demonstrating significant improvement over the baseline. We also analyze the impact of prompt format design and detection thresholds, and provide a qualitative analysis of success and failure cases. Our modular, config-driven pipeline enables reproducible experimentation and serves as a foundation for further extensions.

Index Terms—Open-vocabulary object detection, Grounding DINO, COCO, fine-tuning, vision-language models

I. INTRODUCTION

Traditional object detection systems are constrained to a fixed set of categories defined during training. Adding new categories requires collecting annotated data and retraining or fine-tuning the model — a process that is both expensive and time-consuming. Open-vocabulary object detection aims to overcome this limitation by detecting objects described by arbitrary text prompts, leveraging the joint understanding of visual and language modalities. This capability has transformative potential across diverse real-world applications: robotic manipulation systems that must interact with novel objects in unstructured environments; medical image analysis for detecting rare pathologies not represented in training data; autonomous driving systems encountering unusual roadside objects; and large-scale image retrieval for e-commerce catalogs with dynamically evolving product taxonomies.

Formally, given an input image $I \in \mathbb{R}^{H \times W \times 3}$ and a text query T describing a set of categories $C_T = \{c_1, c_2, \dots, c_K\}$, an open-vocabulary detector must produce a set of N detections $\{(b_i, c_i, s_i)\}_{i=1}^N$, where $b_i = (x_i, y_i, w_i, h_i)$ denotes a bounding box in COCO coordinate format, $c_i \in C_T$ is the predicted category label, and $s_i \in [0, 1]$ is the detection confidence score. Unlike traditional detectors where C_T is a fixed set known at training time, open-vocabulary detection requires the model to generalize to arbitrary C_T specified at inference time.

This task presents three core challenges rooted in the visual-language interface:

- **Visual-language alignment:** The model must understand both image content and text descriptions, mapping semantic concepts to visual regions. This requires joint embedding spaces where textual descriptions of “person” or “bicycle” activate corresponding visual feature regions.
- **Generalization:** The model must detect categories unseen during training, requiring robust visual-language representations that do not overfit to training category distributions.
- **Fine-grained grounding:** The model must localize specific objects described by natural language phrases with precise bounding boxes. This demands token-level alignment between text phrases and image regions.

This project reproduces the Grounding DINO [1] pipeline and further fine-tunes it on the COCO 2017 training set. Our specific contributions are: (1) a complete zero-shot evaluation on the COCO val2017 benchmark (5,000 images) achieving an mAP of 0.485 using the Swin-Tiny backbone; (2) fine-tuning the model for 6 epochs, with epoch 5 yielding the best performance (mAP=0.559, AP50=0.735, AP75=0.615); (3) an analysis of the prompt format, confirming the critical importance of dot-separated category enumeration; (4) a structured error taxonomy identifying small-object detection and crowded-scene confusion as primary limitations; and (5) an open-source, config-driven reproduction pipeline built on the official `groundingdino-py` package.

The remainder of this report is organized as follows. Section 2 reviews related work on Grounding DINO and open-vocabulary detection. Section 3 describes our reproduction methodology, including model configuration, inference pipeline, and prompt design. Section 4 presents experimental results, including baseline evaluation, fine-tuning results, and ablation studies. Section 5 discusses limitations and future work. Section 6 concludes with key findings.

II. RELATED WORKS

A. DETR and DINO: Foundations

Grounding DINO builds upon the DETR (DEtection TRansformer) family of end-to-end object detectors. DETR [2] reformulated object detection as a set prediction problem using a Transformer encoder-decoder architecture with learnable

object queries, eliminating the need for hand-crafted components such as anchor boxes and non-maximum suppression. However, DETR suffered from slow training convergence and difficulty detecting small objects. Deformable DETR [3] addressed these limitations by introducing deformable attention modules that attend to sparse spatial locations around reference points, achieving faster convergence and improved multi-scale feature handling.

DINO [4] (DETR with Improved denoising anchor boxes) further advanced the DETR paradigm through three key innovations: contrastive denoising training, which accelerates convergence by training the model to reconstruct ground-truth objects from noised queries; mixed query selection, which initializes object queries from encoder features rather than learned embeddings; and a look-forward-twice scheme for box prediction refinement. These improvements made DINO competitive with classical detectors while retaining the end-to-end set prediction framework. Grounding DINO [1] extends DINO by fusing text features from a BERT language encoder into the DINO detection architecture, enabling the object queries to become text-conditioned for open-vocabulary detection.

B. Grounding DINO

Grounding DINO [1] extends the DINO detection architecture with grounded pre-training for open-set detection. The key architectural components include:

- **Swin Transformer backbone** [5]: Provides hierarchical visual features through four stages with decreasing spatial resolution ($H/4 \times W/4$ to $H/32 \times W/32$) and increasing channel dimension (96 to 768). The shifted-window attention mechanism enables cross-window information exchange while maintaining linear computational complexity with respect to image size.
- **BERT language encoder** [6]: Encodes free-form text prompts into contextualized token embeddings. We use the bert-base-uncased variant (12 Transformer layers, hidden size 768, 12 attention heads) with a maximum text length of 256 tokens.
- **Deformable attention fusion**: Six cross-modal feature fusion layers, each with 8 attention heads and 256-dimensional features, align visual regions with text phrases through deformable attention.

The model uses 900 object queries that interact with both visual features and text features through the fusion layers. Each query predicts a bounding box, a confidence score, and a text-phrase association. The model is pre-trained on diverse detection and grounding datasets including Objects365, OpenImages, and GoldG.

C. Related Open-Vocabulary Detectors

Open-vocabulary object detection has seen rapid progress since the introduction of vision-language pre-training. Table I compares representative methods.

TABLE I
COMPARISON OF OPEN-VOCABULARY OBJECT DETECTORS

Method	Backbone	Language	Key Innovation	Year
GLIP [7]	Swin-T	BERT	Unified detection + grounding	2022
OWL-ViT [8]	ViT-B/32	CLIP	Vision Transformer + CLIP	2022
Grounding DINO [1]	Swin-T	BERT	DINO + grounded pre-training	2024
YOLO-World [9]	YOLOv8	CLIP	Real-time re-parameterization	2024

Among these, Grounding DINO is chosen for our reproduction due to its strong zero-shot performance, well-maintained open-source implementation, and architectural clarity.

D. COCO Benchmark

Microsoft COCO (Common Objects in Context) [10] is the standard benchmark for object detection evaluation. The dataset contains 80 object categories spanning common everyday objects. The 2017 validation set contains 5,000 images with 36,781 ground-truth bounding box annotations. Evaluation uses the standard COCO metrics: average precision (AP) averaged over 10 IoU thresholds from 0.50 to 0.95, along with per-size breakdowns (APS, APM, APL).

III. METHOD

We reproduce the Grounding DINO pipeline using the official `groundingdino-py` package [1] and extend it with fine-tuning capabilities. Our work focuses on wrapping the official implementation with a modular Python interface for configurable inference, evaluation, and training.

A. Model Configuration

Table II summarizes the model specification used in our reproduction.

TABLE II
COMPLETE GROUNDING DINO MODEL SPECIFICATION

Component	Parameter	Value
Swin-T Backbone [5]	window_size / embed_dim depths / num_heads	7 / 96 [2,2,6,2] / [3,6,12,24]
BERT Encoder [6]	model / hidden_size max_text_len / vocab_size	bert-base-uncased / 768 256 / 30,522
Deformable Fusion	d_model / nhead / num_layers dim_feedforward	256 / 8 / 6 1024
Detection Head	num_queries / num_classes	900 / 80
Input Preprocessing	min_size / max_size mean / std (ImageNet)	800 / 1333 [0.485,0.456,0.406] / [0.229,0.224,0.225]
Inference	box_threshold / text_threshold	0.35 / 0.25 (default)

B. Inference Pipeline




Figure 3: The framework of Grounding DINO. We present the overall framework, a feature enhancer layer, and a decoder layer in block 1, block 2, and block 3, respectively.

Fig. 1. Grounding DINO architecture: Swin Transformer backbone extracts hierarchical visual features; BERT encoder produces contextualized text embeddings; deformable attention fusion aligns visual regions with text tokens through 6 cross-modal layers with 8 attention heads each. 900 object queries jointly attend to both modalities.

The inference pipeline consists of four stages:

Stage 1 — Model Loading. The `GroundingDINOModel` class wraps the official `groundingdino-py` package, loading a 662MB Swin-T checkpoint and validating configuration.

Stage 2 — Image Preprocessing. Input images undergo BGR→RGB conversion, resizing (short side 800px, long side ≤ 1333 px preserving aspect ratio), and ImageNet normalization (mean=[0.485,0.456,0.406], std=[0.229,0.224,0.225]).

Stage 3 — Text Encoding and Cross-Modal Fusion. The dot-separated prompt is tokenized by BERT’s WordPiece tokenizer (vocabulary 30,522 tokens, max length 256). Multi-scale visual features from the Swin-T backbone and contextual text embeddings from BERT are fused through 6 deformable attention layers with 8 heads and 256-dimensional features. Each of the 900 object queries predicts a bounding box, confidence score, and text-phrase association.

Stage 4 — Postprocessing. Normalized (cx, cy, w, h) coordinates are converted to pixel coordinates, filtered by box and text confidence thresholds, and mapped to COCO category IDs. Boxes are clamped to image boundaries and degenerate boxes (width or height ≤ 0) are discarded.

Inference Speed. On an RTX 4060 GPU, single-image inference takes approximately 0.81 seconds. Full COCO val2017 evaluation (5,000 images) completes in approximately 68 minutes. GPU memory usage peaks at approximately 3.2 GB of the available 8 GB VRAM.

C. Prompt Design and Tokenizer Analysis

The text prompt uses a dot-separated format enumerating all 80 COCO categories (e.g., “person . bicycle . car . motorcycle . airplane”). This format is not merely a convention but a functional requirement rooted in BERT’s WordPiece tokenization.

Why Dot-Separated? In the dot-separated format, the period surrounded by spaces is tokenized as an independent separator token, clearly delineating category boundaries. Each

category name is tokenized independently. In contrast, comma-separated format (“person, bicycle, car,”) causes the comma to merge with the preceding word into a single ambiguous token (e.g., “car,”), confusing the model’s phrase-to-category association. Sentence-style formats introduce function words (“There are”, “and”, “in the image”) as noise tokens that dilute attention to category names.

Phrase-to-Category Mapping. After the model outputs detected phrases, a three-layer matching strategy maps each phrase to a COCO category ID:

- 1) **Exact match:** Case-insensitive comparison with the 80 COCO category names.
- 2) **Substring match:** If the phrase partially matches a longer or shorter category name (e.g., “dining” → “dining table”).
- 3) **Alias match:** A manually curated dictionary of 21 common aliases (e.g., “sofa” → “couch”, “tv” → “tv”).

Phrases failing all three layers (approximately 2–5% of detections) are discarded, reducing effective recall.

D. Fine-Tuning Setup

We fine-tune the pre-trained Grounding DINO Swin-T checkpoint on the COCO train2017 set (118K images) for 6 epochs. The training configuration is as follows:

- Optimizer: AdamW with learning rate 5×10^{-5}
- Batch size: 2 (due to GPU memory constraints)
- Mixed precision (AMP) enabled
- Checkpoint saved after every epoch
- Validation performed after every epoch on COCO val2017

We use the same prompt format and threshold settings as in inference. The fine-tuning is implemented using MMDetection and MMEEngine frameworks.

E. Software Architecture

The reproduction is implemented as a config-driven Python package (`cv-project-2026`) consisting of five source modules and ten CLI scripts, totaling approximately 1,826 lines of code. Table III summarizes the module organization.

TABLE III
PROJECT MODULE ORGANIZATION

Module	Responsibility	LOC
<code>src/models/grounding_dino.py</code>	Official model wrapper	182
<code>src/inference/predictor.py</code>	High-level inference API	229
<code>src/inference/visualizer.py</code>	Detection box drawing	262
<code>src/inference/coco_visualizer.py</code>	COCO GT+pred comparison	213
<code>src/engine/evaluator.py</code>	COCOeval integration	368
<code>src/datasets/coco_categories.py</code>	Category mapping & prompts	269
<code>src/datasets/coco_dataset.py</code>	Dataset utilities	188
<code>src/utils/box_ops.py</code>	IoU, NMS, coordinate ops	115
Total		1,826

All configuration is managed through OmegaConf YAML files, enabling CLI overrides for paths and hyperparameters. The architecture follows three design principles: (1) **config-driven** — all settings in YAML with dot-notation CLI overrides; (2) **modular** — each

module has a single responsibility with well-defined interfaces; (3) **error-tolerant** — per-image try/except handling prevents single-image failures from interrupting batch evaluation. Ten CLI scripts support the full workflow: `inference.py`, `eval.py`, `run_experiments.py`, `download_weights.py`, `download_coco.py`, `generate_report_highlights.py`, `generate_visualizations.py`, `train.py`, and platform-specific setup scripts.

IV. EXPERIMENTS

A. Datasets

We evaluate on the COCO 2017 validation set [10], which contains 5,000 images spanning 80 object categories with 36,781 ground-truth bounding box annotations. For fine-tuning, we use the COCO train2017 set (118K images). All experiments use the standard COCO evaluation metrics.

B. Implementation Details

All experiments were conducted on an NVIDIA RTX 4090 GPU for training and an RTX 4060 for inference. The software environment includes Python 3.10, PyTorch 2.7.1, CUDA 11.8, and the `groundingdino-py` package (v0.4.0). Default hyperparameters for inference are `box_threshold=0.35` and `text_threshold=0.25`.

C. Metrics

We use the standard COCO evaluation metrics: AP (averaged over IoU thresholds 0.50:0.95), AP50, AP75, APS, APM, APL, as well as AR@1, AR@10, AR@100 with per-size breakdowns.

D. Experimental Design

We first evaluate the official pre-trained Grounding DINO Swin-T checkpoint as the baseline (zero-shot). Then, we fine-tune the model on the COCO training set for 6 epochs, saving a checkpoint after each epoch. Each checkpoint is evaluated on the COCO validation set. The goal is to observe whether fine-tuning improves detection performance and to select the best checkpoint for the final demo.

E. Main Results

Table IV presents the complete evaluation results for the baseline and each fine-tuned epoch.

TABLE IV
COCO VAL2017 EVALUATION RESULTS (BASELINE AND FINE-TUNED CHECKPOINTS)

Model	mAP	AP50	AP75	APs	APm	APl
Baseline	0.485	0.645	0.530	0.342	0.518	0.635
Epoch 1	0.546	0.718	0.600	0.379	0.583	0.693
Epoch 2	0.553	0.722	0.606	0.400	0.592	0.701
Epoch 3	0.556	0.726	0.610	0.393	0.594	0.703
Epoch 4	0.555	0.729	0.611	0.394	0.595	0.703
Epoch 5	0.559	0.735	0.615	0.400	0.593	0.712
Epoch 6	0.558	0.733	0.614	0.394	0.590	0.704

Epoch 5 achieves the best overall performance (mAP=0.559) and is selected as the final checkpoint for the demo. Fine-tuning significantly improves the model compared with the pretrained baseline: overall mAP improves from 0.485 to 0.559, an absolute gain of 0.074. AP50 improves from 0.645 to 0.735, showing better object recognition and coarse localization. AP75 improves from 0.530 to 0.615, indicating more accurate bounding box localization. The model nearly converges after 5 epochs, as epoch 6 shows a slight drop.

F. Result Analysis




Fig. 2. Full COCO metrics overview: AP, AP50, AP75, and per-size breakdown (APS, APM, APL) for the best checkpoint (epoch 5).

The per-size metrics reveal a clear size-dependent performance gradient. Large objects (APL=0.712) achieve approximately $1.8\times$ the performance of small objects (APS=0.400), consistent with the detection literature. This gap is driven by the Swin-T backbone’s effective resolution: at stride 32 in the deepest layer, small objects ($<32\times32$ pixels) occupy very few feature map locations, limiting the model’s ability to both classify and precisely localize them.




Fig. 3. AP gap across object sizes: small (APS), medium (APM), and large (APL) objects.

Fine-tuning significantly improves all metrics: overall mAP from 0.485 (baseline) to 0.559 (epoch 5), an absolute gain of 0.074. AP50 improves from 0.645 to 0.735, demonstrating better object recognition and coarse localization. AP75 improves from 0.530 to 0.615, indicating more accurate bounding box regression. The model nearly converges after 5 epochs, as epoch 6 shows negligible improvement or slight regression.

Qualitative analysis of detection outputs is shown in Figures 4 and 5. The model successfully detects most medium and large objects with correct category labels and tight bounding boxes across diverse scenes including indoor living spaces, outdoor street views, and crowded public areas. Detection confidence scores for clearly visible objects typically exceed 0.5.




Fig. 4. Success cases: accurate detections with correct categories and tight bounding boxes for medium and large objects across diverse scene types.

Common failure modes include:

- **Missed small objects (most frequent):** Objects smaller than 32×32 pixels are frequently undetected, accounting for the majority of false negatives. Distant pedestrians, small traffic signs, and kitchen utensils in wide-angle scenes are typical examples.
- **Crowded-scene confusion (occasional):** In dense scenes with many overlapping instances (> 15 ground-truth objects), the model may produce duplicate bounding boxes or mislabel partially occluded instances.
- **Occlusion-related localization errors (moderate):** Partial occlusion reduces detection confidence and degrades bounding box precision. This explains part of the gap between AP50 and AP75: the model can roughly locate occluded objects but struggles to produce tight bounding boxes.




Fig. 5. Failure cases: small object misses (top), crowded-scene confusion (middle), and occlusion-related errors (bottom).

G. Prompt Format Analysis

We verify the importance of prompt format through controlled experiments. Three formats were compared on a 500-image subset: dot-separated (“person . bicycle . car .”), comma-separated (“person, bicycle, car,”), and sentence-style (“There are person, bicycle, car...”). Table V presents the results.

TABLE V
PROMPT FORMAT COMPARISON (SUBSET 500 IMAGES)

Format	AP	AP50	Boxes/Img
dot-separated	0.4382	0.5532	7.72
comma-separated	0.1671	0.2038	2.61
sentence-style	0.1675	0.2051	2.49

The dot-separated format dramatically outperforms alternatives across all metrics, with approximately 62% higher AP. The performance gap is most severe for small objects. This sensitivity stems from BERT’s WordPiece tokenizer behavior: the period surrounded by spaces is tokenized as a distinct separator, while commas merge with adjacent words into ambiguous tokens. Sentence-style formats add function words that dilute attention to category names. This finding has direct implications for practical deployment: any system using Grounding DINO must format text prompts with dot separators.




Fig. 6. Prompt format comparison: dot-separated, comma-separated, and sentence-style.

H. Threshold Sensitivity

We also conduct a threshold sensitivity analysis on the baseline zero-shot model. Table VI shows results across three threshold settings on a 500-image subset.

TABLE VI
THRESHOLD SENSITIVITY (SUBSET 500 IMAGES)

box_th	text_th	AP	AP50	AP75	Boxes/Img
0.25	0.20	0.4637	0.5884	0.5083	13.85
0.35	0.25	0.4382	0.5532	0.4793	7.72
0.45	0.30	0.3931	0.4827	0.4317	5.08

Lower thresholds yield higher AP but at the cost of significantly more detections. Dropping box_threshold from 0.35 to 0.25 increases the average detection count by 79% (7.72 to 13.85 per image) while improving AP by 5.8%. Small objects benefit most from lower thresholds: APS increases by 46% when lowering box_th from 0.45 to 0.25, compared to only 8.5% improvement for APL. This indicates that small-object detections tend to have lower confidence scores and are more prone to being pruned by aggressive thresholding. For our fine-tuned model, we retain the default thresholds (box_th=0.35, text_th=0.25) for consistency.




Fig. 7. Threshold sensitivity analysis showing AP variation across box and text threshold combinations.

V. DISCUSSION

A. Interactive Demo Interface

As a supplementary component of this project, we developed an interactive web-based demonstration interface for real-time Grounding DINO inference. The interface provides the following capabilities:

- **Live image upload and detection:** Users can upload arbitrary images and receive detection results within seconds, with annotated bounding boxes, category labels, and confidence scores overlaid.
- **Custom text prompts:** Users can input arbitrary text prompts (not limited to COCO categories), enabling open-vocabulary detection of any object describable in natural language.
- **Adjustable detection parameters:** Interactive sliders allow real-time adjustment of box_threshold and text_threshold, enabling users to explore the precision-recall tradeoff interactively.
- **Multi-image batch processing:** The interface supports uploading multiple images for batch detection, with results displayed in a gallery format for side-by-side comparison.
- **Result export:** Annotated images and detection results (bounding boxes, scores, phrases) can be downloaded in JSON format for further analysis.

The demo is built as a lightweight wrapper around the Predictor class, demonstrating the modularity of the project's software architecture. The same predictor instance used for COCO evaluation serves inference requests from the web interface with no code duplication.

B. Comparison with Reported Benchmarks

The official Grounding DINO paper reports a zero-shot AP of approximately 0.485 for the Swin-T backbone, which matches our baseline. After fine-tuning on COCO, the paper reports an AP of approximately 0.506. Our fine-tuned result (AP=0.559) exceeds this reported value, likely due to differences in training hyperparameters, longer training (6 epochs vs. perhaps fewer), or implementation details. This demonstrates that our reproduction is successful and even yields improved performance.

C. Limitations

The primary limitation remains small-object detection (APS=0.400 even after fine-tuning), attributable to the Swin-T backbone's resolution constraints. The computational cost of fine-tuning (6 epochs on RTX 4090) is non-negligible, but the resulting model achieves strong performance. The prompt format brittleness is another practical limitation for end-user applications.

D. Future Work

Future work should explore multi-scale feature fusion, higher input resolution, or alternative backbones to improve small-object detection. Additionally, cross-dataset evaluation

and deployment optimization (e.g., TensorRT) are promising directions. The modular pipeline developed in this project facilitates these extensions.

VI. CONCLUSION

This reproduction successfully demonstrates Grounding DINO’s open-vocabulary detection capability and further improves it via fine-tuning on COCO. Our key findings are:

- The pre-trained model achieves a zero-shot mAP of 0.485 on COCO val2017.
- Fine-tuning for 5 epochs yields the best performance: mAP=0.559, AP50=0.735, AP75=0.615.
- Small-object detection remains the main bottleneck.
- Dot-separated prompts are essential for the model’s text encoder.

The modular, config-driven pipeline provides a solid foundation for future research and deployment. All code and checkpoints are available in our public repository for reproducibility.

REFERENCES

- [1] S. Liu, Z. Zeng, T. Ren, F. Li, H. Zhang, J. Yang, Q. Jiang, C. Li, J. Yang, H. Su, J. Zhu, and L. Zhang, “Grounding DINO: Marrying DINO with grounded pre-training for open-set object detection,” in *European Conference on Computer Vision (ECCV)*, 2024.
- [2] N. Carion, F. Massa, G. Synnaeve, N. Usunier, A. Kirillov, and S. Zagoruyko, “End-to-end object detection with transformers,” in *European Conference on Computer Vision (ECCV)*, 2020.
- [3] X. Zhu, W. Su, L. Lu, B. Li, X. Wang, and J. Dai, “Deformable DETR: Deformable transformers for end-to-end object detection,” in *International Conference on Learning Representations (ICLR)*, 2021.
- [4] H. Zhang, F. Li, S. Liu, L. Zhang, H. Su, J. Zhu, L. M. Ni, and H.-Y. Shum, “DINO: DETR with improved denoising anchor boxes for end-to-end object detection,” in *International Conference on Learning Representations (ICLR)*, 2023.
- [5] Z. Liu, Y. Lin, Y. Cao, H. Hu, Y. Wei, Z. Zhang, S. Lin, and B. Guo, “Swin transformer: Hierarchical vision transformer using shifted windows,” in *IEEE/CVF International Conference on Computer Vision (ICCV)*, 2021.
- [6] J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “BERT: Pre-training of deep bidirectional transformers for language understanding,” in *Conference of the North American Chapter of the Association for Computational Linguistics (NAACL)*, 2019.
- [7] L. H. Li, P. Zhang, H. Zhang, J. Yang, C. Li, Y. Zhong, L. Wang, L. Yuan, L. Zhang, J.-N. Hwang, K.-W. Chang, and J. Gao, “Grounded language-image pre-training,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.
- [8] M. Minderer, A. Gritsenko, A. Stone, M. Neumann, D. Weissenborn, A. Dosovitskiy, A. Mahendran, A. Arnab, M. Dehghani, Z. Shen, X. Wang, X. Zhai, T. Kipf, and N. Houlsby, “Simple open-vocabulary object detection with vision transformers,” in *European Conference on Computer Vision (ECCV)*, 2022.
- [9] T. Cheng, L. Song, Y. Ge, W. Liu, X. Wang, and Y. Shan, “YOLO-world: Real-time open-vocabulary object detection,” in *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2024.
- [10] T.-Y. Lin, M. Maire, S. Belongie, J. Hays, P. Perona, D. Ramanan, P. Dollár, and C. L. Zitnick, “Microsoft COCO: Common objects in context,” in *European Conference on Computer Vision (ECCV)*, 2014.

APPENDIX A REPRODUCTION COMMANDS

The following commands reproduce all experiments described in this report. All paths are relative to the repository root.

A. Environment Setup

```
# Create and activate conda environment
conda create -n grounding_dino python=3.10 -y
conda activate grounding_dino

# Install PyTorch with CUDA 11.8
pip install torch torchvision --index-url https://download.pytorch.org/whl/cu118

# Install core dependencies
pip install groundingdino-py
pip install -r requirements.txt
pip install -e ".[dev]"
```

B. Data and Model Download

```
# Download Grounding DINO checkpoint (662 MB)
python scripts/download_weights.py

# Download COCO 2017 val2017 and annotations
python scripts/download_coco.py
```

C. Single-Image Inference

```
python scripts/inference.py \
    --image data/demo_images/000000000139.jpg \
    --text "person . bicycle . car . motorcycle . bus ." \
    --checkpoint checkpoints/groundingdino_swint_ogc.pth
```

D. Full COCO val2017 Evaluation

```
python scripts/eval.py \
    --config configs/grounding_dino.yaml \
    --checkpoint checkpoints/groundingdino_swint_ogc.pth \
    --coco_image_dir data/coco/val2017 \
    --coco_ann_file data/coco/annotations/instances_val2017.json \
    --output_dir outputs/coco_eval/full_val2017
```

E. Fine-Tuning (6 Epochs)

```
python scripts/train.py \
    --config configs/grounding_dino_train.yaml \
    --checkpoint checkpoints/groundingdino_swint_ogc.pth \
    --train_ann_file data/coco/annotations/instances_train2017.json \
    --val_ann_file data/coco/annotations/instances_val2017.json \
    --output_dir outputs/finetune
```

APPENDIX B COMPLETE COCO PROMPT

The full dot-separated prompt used for COCO evaluation, enumerating all 80 categories:

```

person . bicycle . car . motorcycle . airplane . bus
. train . truck . boat . traffic light . fire
hydrant . stop sign . parking meter . bench .
bird . cat . dog . horse . sheep . cow .
elephant . bear . zebra . giraffe . backpack .
umbrella . handbag . tie . suitcase . frisbee .
skis . snowboard . sports ball . kite . baseball
bat . baseball glove . skateboard . surfboard .
tennis racket . bottle . wine glass . cup .
fork . knife . spoon . bowl . banana . apple .
sandwich . orange . broccoli . carrot . hot dog
. pizza . donut . cake . chair . couch . potted
plant . bed . dining table . toilet . tv .
laptop . mouse . remote . keyboard . cell phone
. microwave . oven . toaster . sink .
refrigerator . book . clock . vase . scissors .
teddy bear . hair drier . toothbrush .

```

Assisted in the quantitative analysis of experimental results, including the computation of per-size AP breakdowns and detection count statistics. Participated in result visualization by helping generate the report highlight figures and organizing the success/failure case grids. Contributed to the writing of the Experiments section and the Discussion section of this report.

APPENDIX C

OUTPUT FILE STRUCTURE

All evaluation and experiment outputs are organized under outputs/:

```

outputs/
  coco_eval/
    full_val2017/          # Main evaluation (5000
    images)
    predictions.json       # Detection results in
    COCO format
    metrics.json           # Full 12 AP+AR metrics
    eval.log               # Detailed evaluation log
    phrase_mapping.json    # Phrase-to-category
    mapping
  finetune/                # Fine-tuned checkpoints
  and evaluations
    epoch_1/ ... epoch_6/
  experiments/
    threshold_sensitivity/ # Optional ablation
    prompt_comparison/     # Optional ablation
  visualizations/
    report_highlights/     # Report figures
    coco_eval/            # Success/failure case
    images
    inference_demo/        # Single-image inference
    examples
  inference_results/       # Batch inference outputs

```

CONTRIBUTIONS

Su Siqi (12310645): Contributed to the literature review and background research on open-vocabulary object detection methods. Assisted in the design and execution of the prompt format ablation experiments. Participated in the qualitative error analysis by manually reviewing detection outputs to identify and categorize failure modes.

Yang Guanyuhan (12313614): Responsible for the overall project design and implementation. Contributions include: wrapping the official groundingdino-py inference backend with a modular Python interface, building the high-level predictor and visualizer, implementing the COCO evaluator with pycocotools integration, developing dataset utilities for category mapping and prompt construction, and implementing box operation utilities. Created all CLI scripts for inference, evaluation, ablation experiments, and visualization generation. Conducted the full COCO val2017 evaluation and all fine-tuning experiments. Authored the project report.

Mo Fengyuan (12311805): Contributed to the design and execution of the threshold sensitivity ablation experiments.